



U.S. General Services
Administration




Chapter 3.1B: Implement Evaluation Pipelines -- Guided Practice

US General Service Administration
AI Community of Practice - March 2026

Why Generic Metrics Are Not Enough

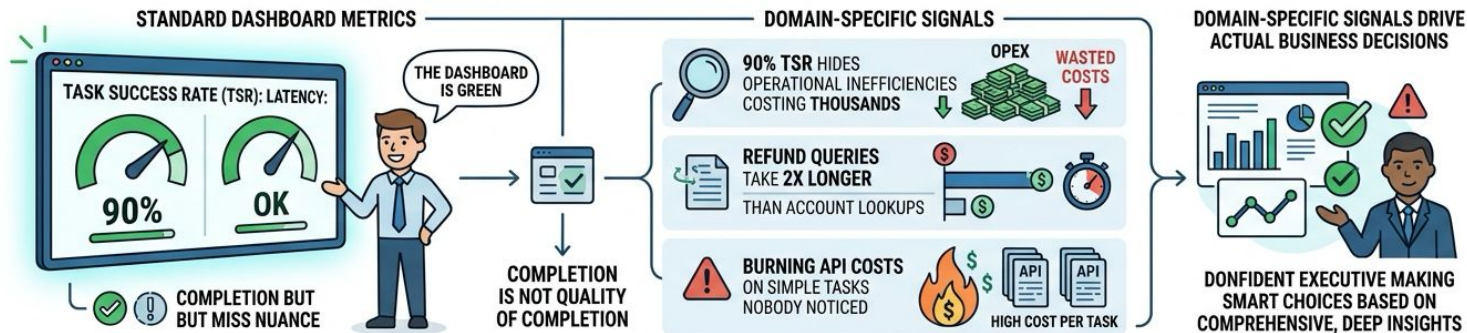
ADAPTING EVALUATION SYSTEMS | From Observation to Application



MANUAL TESTING UNDER DEADLINE PRESSURE IS SKIPPED

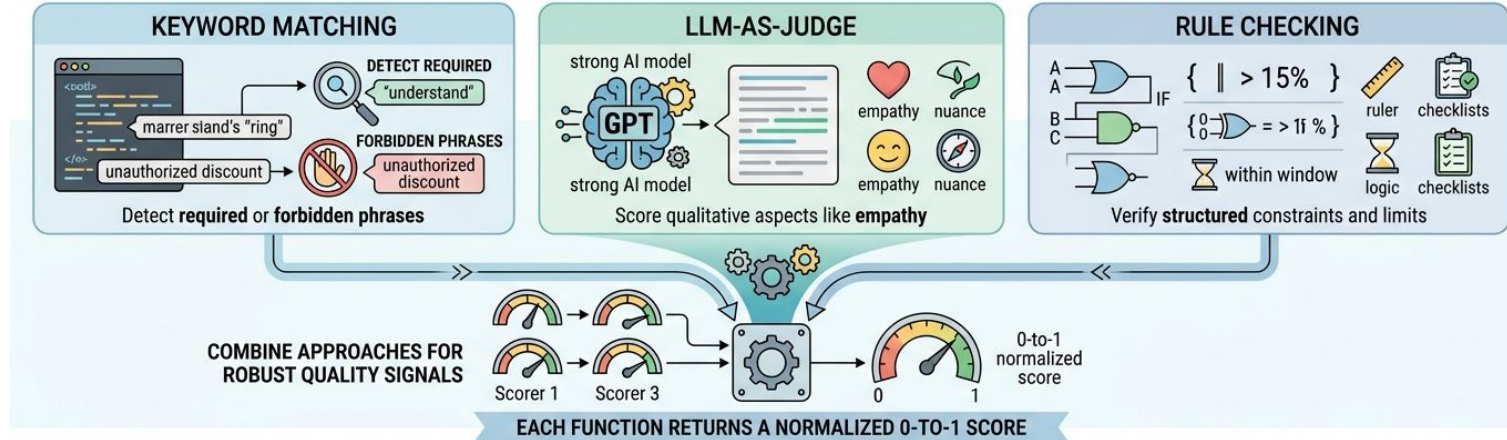


WHY GENERIC METRICS ARE NOT ENOUGH

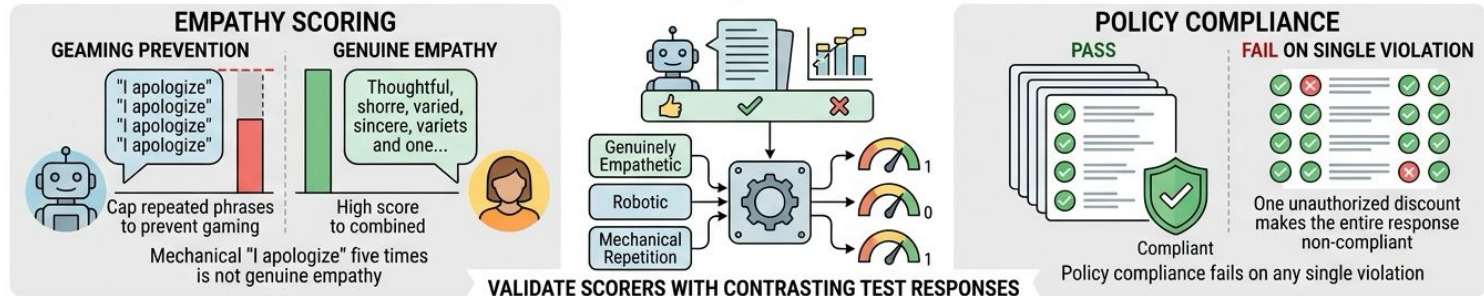


Custom Scoring & Metrics

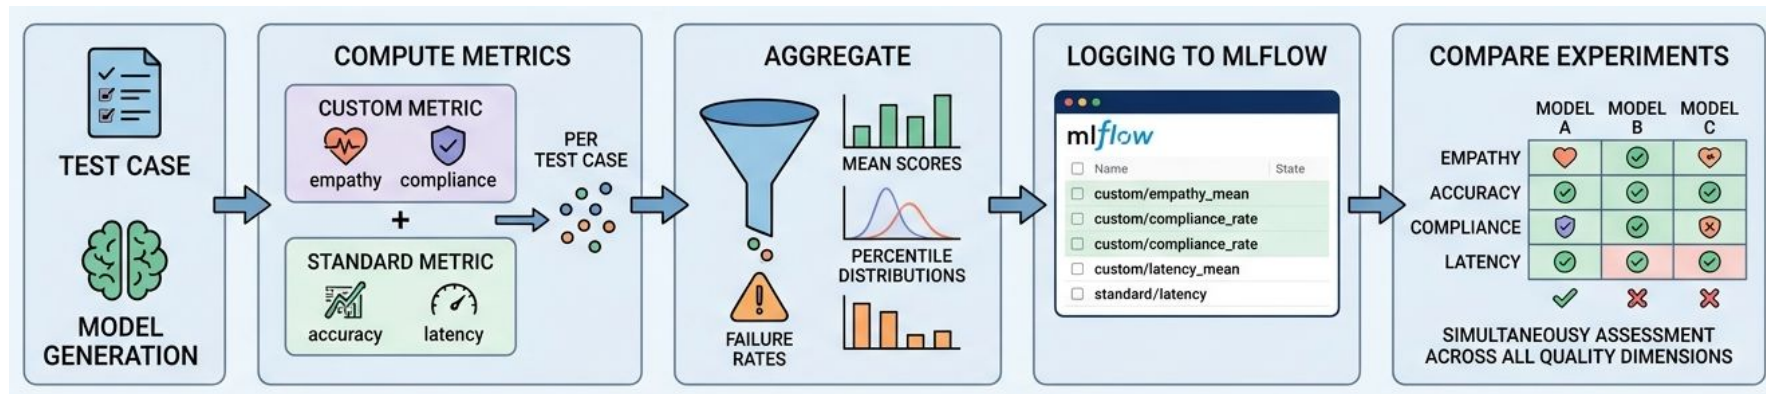
THREE APPROACHES TO CUSTOM SCORING



EMPATHY AND COMPLIANCE – METRICS THAT MATTER



Integrating Custom Metrics into MLflow



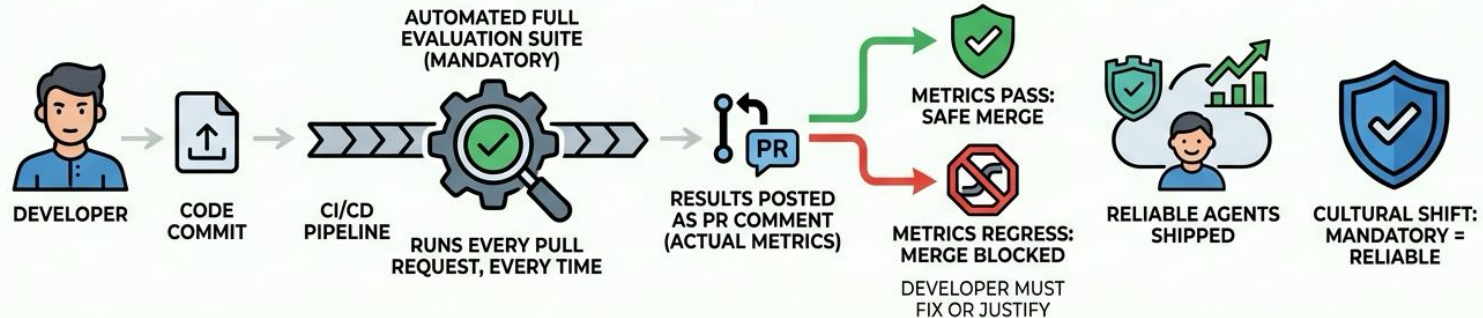
- Compute custom metrics per test case alongside standard ones
- Do not defer custom scoring to a separate pass
- Aggregate: mean scores, percentile distributions, failure rates
- Log with consistent naming prefixes like "custom/"
- Compare experiments across all quality dimensions at once

Why CI/CD Evaluation Is Non-Negotiable

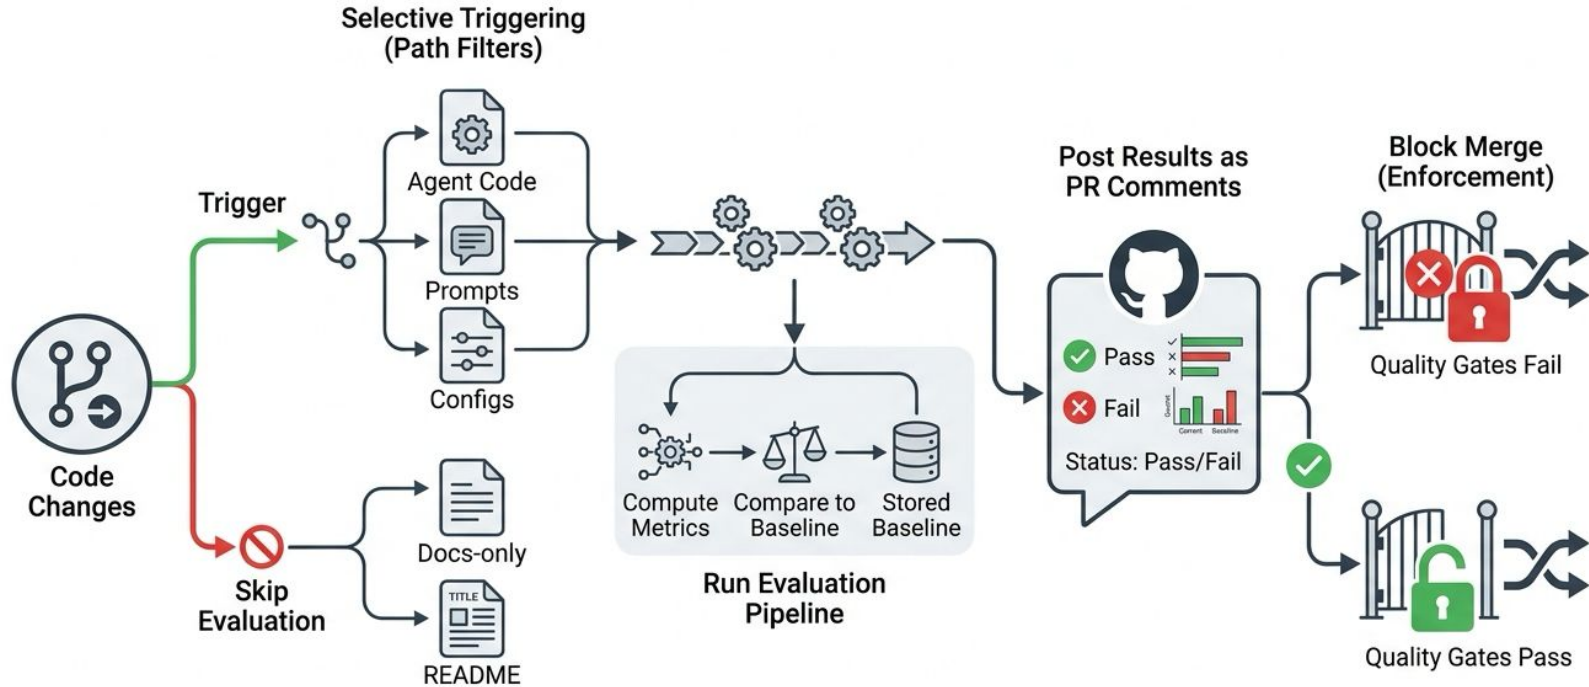
MANUAL EVALUATION: SKIPPED UNDER PRESSURE (RISK)



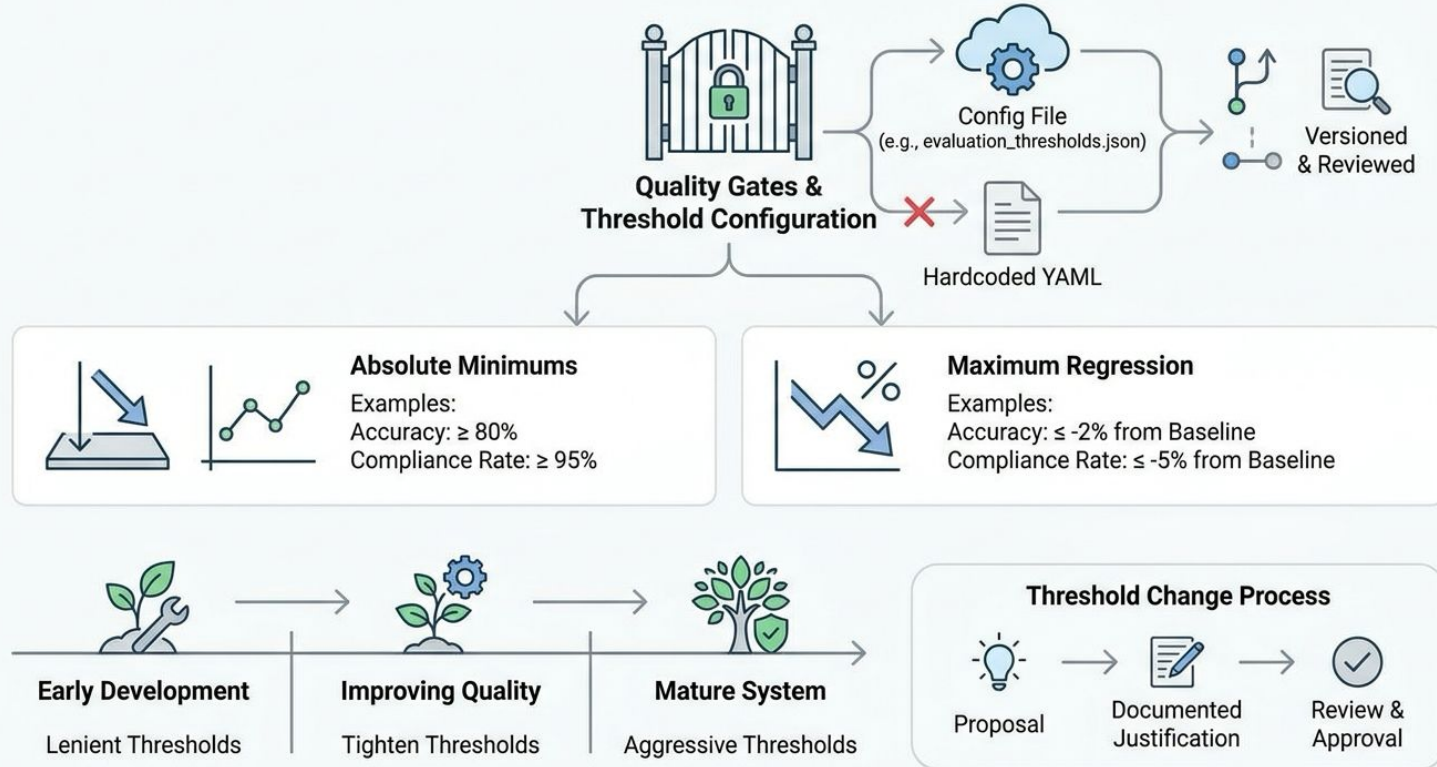
AUTOMATED CI/CD EVALUATION: MANDATORY (RELIABILITY)




GitHub Actions Evaluation Workflow




Quality Gates and Threshold Configuration



Validating Your Evaluation Pipeline




Closing the Feedback Loop



Key Takeaways

KEY PRINCIPLES FOR ROBUST DEPLOYMENT

GO BEYOND GENERIC METRICS



Measure specific domain needs

THREE SCORING APPROACHES

KEYWORDS



Detect exact terms

LLM-AS-JUDGE



Evaluate qualitative aspects

MANDATORY QUALITY GATES

RULES



Enforce defined constraints


MANDATORY QUALITY GATES



MANDATORY QUALITY GATES

Prevent regressions from production

PROGRESSIVE THRESHOLDS



PROGRESSIVE THRESHOLDS

Configure thresholds across lifecycle

VALIDATE PIPELINE WITH PRs



VALIDATE PIPELINE WITH PRs

Validate pipeline behavior